


NAVIGATING THE

ETHICS OF AI

MORALITY, LAW AND INNOVATION IN THE AGE OF AI





Alva Markelius
Guest Lecture
KTH 21st of September 2026

INTRODUCTION

Alva Markelius

- PhD candidate, Cambridge Trust Scholar, Affective Intelligence and Robotics Laboratory (AFAR), Department of Computer Science, University of Cambridge
- Top 100 Brilliant Women in AI Ethics™ 2024
- MSt AI Ethics & Society (Leverhulme Centre for the Future of Intelligence, University of Cambridge)
- BSc Cognitive Science (University of Gothenburg + Seoul National University)
- Fun fact: I am president of Cambridge Uni Polo Club 🐎




TODAY'S LECTURE

- Part 1: Introduction to AI Ethics (45 minutes)
 - Taxonomy of the field
 - Concepts and important topics
 - Conceptualising AI Ethics
 - The EU AI Act
- Part 2: Current research topics and challenges in AI Ethics (30 mins)
 - AI Hype - potentials and pitfalls
 - Sustainability and planetary impact
 - Social AI and Robotics
- Q&A (15 mins)

PART 1: INTRODUCTION TO AI ETHICS


WHAT IS AI ETHICS?

Trolley problem: the tip of the iceberg!




WHAT IS AI ETHICS?

A Simple Taxonomy of the Field



WHAT IS AI ETHICS?

A Simple Taxonomy of the Field



WHY DOES IT MATTER?

NEWS DEMAND EQUITY

Amazon Shuts Down AI Hiring Tool for Being Sexist


Amazon created an artificial intelligence tool to help with recruiting, but it showed gender bias.

XNyYwVsIERIcGxveXMgRXhwYw
1IEZhY2IhbCBSZWVZ25pdGlw
m9ncmFtIGlReport.3800uIEd

Israel Deploys Expansive Facial Recognition Program in Gaza


nytimes.com · 2024

Within minutes of walking through an Israeli



The first AI-judged beauty contest taught us one thing: Robots are racist

With more than 6,000 applicants from over 100 countries competing, the first international beauty contest judged entirely by artificial intelligence just came to end. The results are a bit disheartening.



I tried the Replika AI companion and can see why users are falling hard. The app raises serious ethical questions

BUSINESS • TECHNOLOGY

Exclusive: OpenAI Used Kenyan Workers on Less Than \$2 Per Hour to Make ChatGPT Less Toxic

Princeton University's 'AI Snake Oil' authors say generative AI hype has 'spiraled out of control'




The Dating App That Knows You Secretly Aren't Into Guys From Other Races

Even if you say "no preference" for ethnicity, the dating app tends to show you people of your own race.

The AI companion who cares

Always here to listen and talk. Always on your side. Join the movement.

US-China AI Race: Policy as Industrial Policy



The environmental pollution behind the boom in artificial intelligence

Huge server farms that power AI have a climate impact that could increase to 14% of global carbon emissions by 2025, a new study says.

Viki Auslender, Shani Ashkenazi 08:37, 23.04.23


TAGS: Environment AI Crypto



Men Are Creating AI Girlfriends and Then Verbally Abusing Them

m.com · 2022

A smartphone app Replika lets users create digital companions, powered by machine learning, that can carry on almost-coherent text conversations.



Students of color are getting flagged to their teachers because testing software can't see them

theverge.com · 2021

Proctorio, a piece of exam surveillance software designed to keep students from cheating while taking tests, relies on open-source software to identify faces.

TECHNOLOGY

Tech leaders urge a pause in the 'out-of-control' artificial intelligence race

March 29, 2023 · 5:05 PM ET

By The Associated Press

Generative AI refers to a class of artificial intelligence (AI) models that can create or generate new data, such as images, text, or music, that is similar to the data it was trained on. Generative models learn to recognize patterns and relationships in the input data and use that knowledge to generate new data.

COMMON TOPICS IN AI ETHICS

Fairness

Inherent biases in machine learning models and algorithms which often arise from skewed training data, algorithmic design or implementation and use-case design, disproportionately affecting certain demographic groups (most often than not reflecting and reinforcing existing structural inequalities).

Accountability

The challenge of determining responsibility when AI systems make errors or harmful decisions. As the AI infrastructure becomes more autonomous and complex, it is increasingly difficult to assign blame or liability for unintended consequences and harms.

Transparency

Clarity and openness with which AI systems, algorithms, and decision-making processes are designed, developed, and deployed. It involves making the inner workings of AI models understandable to users, stakeholders, and regulators, allowing for scrutiny of how data is used, how decisions are made, and what factors influence outcomes.

Explainability

Challenge of making AI systems more understandable/interpretable to humans. This lack of transparency raises concerns about accountability, trust, and the ability to identify and rectify errors or biases in AI-driven outcomes.

a software developer



a flight attendant



a chef



a cook



a taxi driver



a housekeeper



mental Hiring Tool Allegedly Displayed binnings

reportedly developed an AI-powered recruiting
d on a decade of resumes purportedly drawn
ay the system learned to favor male candidates,...

ACCOUNTABILITY

- Displacement of human actors, and introduction of data-driven computerised, algorithmic systems has created **barriers** to accountability:
 - the **extent of involved actors**, especially along machine learning pipelines and usage of multi-purpose, pre-trained toolkits.
 - **failure modes** or performance errors “bugs” in the systems, seen as inevitable and thus may impede the attribution of blame.
 - data-driven algorithmic systems as a **scapegoat for moral agency**, driven by a rhetoric of the systems as intelligent or overreliance on mathematical guarantees
 - **ownership without liability**, complicated by third-party-actors, cyber-physical systems and increased impact of tech companies over regulation.



Judge Rules \$400 Million Algorithmic System Illegally Denied Thousands of People's Medicaid Benefits
Thousands of children and adults were automatically terminated from Medicaid and disability benefits programs by a computer system that was supposed to make applying for and receiving health coverage easier.




TRANSPARENCY


- Deeply intertwined with both accountability and explainability of AI systems - core principle in AI standards and policies, particularly for enabling responsible and ethical AI
- Relates to making the decision making processes and inner workings of **autonomous, opaque and “black-box” AI systems**, especially safety-critical ones, or systems prone to cause harm, more interpretable and understandable to humans
- ALSO includes transparency regarding e.g. **data sources, the rights of the human labour involved, and the documentation of validation processes**
- AI systems are not just collections of code and data; they are complex assemblages of human and nonhuman actors.

EXPLAINABILITY

- How AI systems' processes and outcomes can be understood and communicated by humans
 - how AI systems make decisions, what variables influences these decisions, and the underlying algorithms and models.
- Particularly important in **high-stakes scenarios** where AI decisions can significantly impact individuals' lives, such as in healthcare, finance, and law enforcement
- A lack of explainability, or when AI models act like “black boxes” can lead to opaque decision-making processes, compromise public trust, and hinder compliance with legal and ethical standards



AI ETHICS STRETCHES FAR BEYOND FATE




IS AI...?

AN INFRASTRUCTURE?

AN INSTRUMENT?


AN IDEOLOGY?

AI AS INFRASTRUCTURE




Anatomy of an AI system

An anatomical case study of the Amazon echo as a artificial intelligence system made of human labor




AI AS INFRASTRUCTURE

- Datacentres – energy- and water usage
- Training of large language models (LLMs)
 - Since 2012 the compute needed to train an AI-system has DOUBLED each 3-4th month
- AI is material - hardware and planetary resources



Drought-stricken communities push back against data centers

As cash-strapped cities welcome Big Tech to build hundreds of million-dollar data centers in their backyards, critics question the environmental cost.



The environmental pollution behind the boom in artificial intelligence

Huge server farms that power AI have a climate impact that could increase to 14% of global carbon emissions

Viki Auslender, Shani Ashkenazi 08:37, 23.04.23

TAGS: [Environment](#) [AI](#) [Crypto](#)

from **2024** **Guardian**

AI likely to increase energy use and accelerate climate misinformation – report

Claims that artificial intelligence will help solve the climate crisis are misguided, warns a coalition of environmental groups

the Guardian / Mar 25, 2024

entre industry's

s (LLMs)

ed to train a

3-4th mon



The Uneven Distribution of AI's Environmental Impacts

The training process for a single AI model, such as an LLM, can consume thousands of megawatt hours of electricity and emit hundreds of tons of carbon. AI model training can also lead to the evaporation of an astonishing amount of...

ss Review / Jul 15, 2024

Why water u
dirty little s

GOATS AND SODA

Epic drought in Taiwan pits farmers against high-tech factories for water

April 19, 2023 • 8:38 AM ET



Emily Feng

back

their backyards, critics question the environmental cost.



The environmental pollution behind the boom in artificial intelligence

Huge server farms that power AI have a climate impact that could increase to 14% of global carbon emissions

Viki Auslender, Shani Ashkenazi 08:37, 23.04.23

TAGS: [Environment](#) [AI](#) [Crypto](#)


The computer chip industry has a dirty climate secret

As demand for chips surges, the semiconductor industry is trying to grapple with its huge carbon foot print


AI AS INFRASTRUCTURE

The Cloud Is Material: On the Environmental Impacts of Computation and Data Storage (Monserate, 2022):

- Environmental costs obscured by sheer complexity of infrastructures and supply chains involved
- identifies the various material and ecological impacts of computation and digital data storage practices.
 - carbon emissions, water usage, noise, and toxic e-waste.
- environmental toll of digital life and the complexity of infrastructures involved in its operation - s socio-technical challenge
- By 2026, computing power dedicated to training AI is expected to increase tenfold.
- By 2030, there will be a 160% increase in demand for power propelled by AI applications.




AI AS INFRASTRUCTURE



AI AS IDEOLOGY


Navigating the AI Safety and AI Ethics debate







AI AS IDEOLOGY

Narratives and visualisations

- **AI narratives** have a profound impact on e.g.
 - Dehumanising effect of humanising machines
 - How automation affects society
 - Cultural assumptions existing in anthropomorphisation of AI
 - Importance of understanding how AI is a distributed phenomenon
 - How AI interacts with and defines global geopolitics and nation sovereignty
 - Gendered/racialised expectations
 - Fear of losing control of AI



- **AI visualisations** also has profound ethical consequences:
 - unhelpful clichéd and misleading images
 - hinder understanding of AI's significant societal and environmental impacts
 - set unrealistic expectations
 - mask accountability of human developers
 - potentially sow fear
 - laden with historical assumptions about gender, ethnicity and religion



Clarote & AI4Media / Better Images of AI / Labour/Resources / CC-BY 4.0

Image by Comuzi / © BBC / Better Images of AI / Mirror D / CC-BY 4.0


therine Breslin & Team and Adobe Firefly / Better Images of AI / Chipping Silicon / CC-BY 4.0

Philipp Schmitt / Better Images of AI / Neural network diagram / CC-BY 4.0


Fritzchens Fritz / Better Images of AI / GPU shot etched 5 / CC-BY 4.0

AI AS INSTRUMENT

- The learning algorithm: compressing the world into a statistical model
 - secularize AI from the ideological status of 'intelligent machine' to one of knowledge instrument.
- All models are wrong, but some are useful:
 - notions of data fitting, overfitting, underfitting, interpolation and extrapolation - ultimately proxies
- Ghost work and heteromation
 - The familiar narrative of AI as *perpetuum mobile* is possible only thanks to a reserve army of workers.



THE EU AI ACT



THE ACT REGULATES USES, NOT TECHNOLOGIES

PROVIDER VS. DEPLOYER

THE EU AI ACT - WHAT'S LIVE NOW

- **Feb 2025: Prohibited practices**
 - Social scoring, manipulative techniques, emotion recognition at work and in education, untargeted facial-image scraping.
- **Aug 2025: GPAI (Providers of general-purpose AI) model obligations**
 - Technical documentation, training data summary, copyright policy etc.
- **Aug 2026: Article 50 transparency**
- **Aug 2026: Enforcement powers switch on**

Article 50:

- Chatbots must declare they're AI
- Deepfakes must be labelled
- AI output machine-readable marked
- Emotion recognition disclosed to the person

THE EU AI ACT - WHAT'S COMING

- **2 Dec 2026:** New ban: NCII (Non-Consensual Intimate Imagery) and CSAM (Child Sexual Abuse Material) generators
- **2 Aug 2027:** National sandboxes operational
- **2 Dec 2027:** Annex III high-risk
- **2 Aug 2028:** Annex I high-risk

Annex III — use-based

- Biometrics · employment · education
- Credit · critical infrastructure · migration

Annex I — product-embedded

- Medical devices · lifts · toys · vehicles

THE EU AI ACT - TEETH

Fines (Art. 99)

- €35m / 7% — prohibited practices
- €15m / 3% — most other breaches
- €7.5m / 1% — misleading the regulator
- Higher of the two · SMEs: lower of the two
- % is worldwide turnover

Beyond fines

- Information requests
- Model access for evaluation
- Mandated mitigations
- Withdrawal from the EU market



PART 2: CASE STUDIES AND CURRENT CHALLENGES

AI HYPE - MECHANISMS AND CONSEQUENCES

Mechanisms:



Anthropomorphism



AI “Experts” and exaggerated literacy



FOMO-driven narratives



Over usage of the term "AI" (AI Snake-oil)

AI HYPE - MECHANISMS AND CONSEQUENCES

Mechanisms:



Anthropomorphism



AI “Experts” and exaggerated literacy



FOMO-driven narratives




Over usage of the term "AI" (AI Snake-oil)

Consequences:

- Planetary costs
 - Data centers and water
 - Compute och electricity
 - Precious metals and hardware (chips)
- Social costs
 - Job loss and job polarization
 - Knowledge decay
 - Knowledge corruption and post-truth

AI HYPE - MECHANISMS AND CONSEQUENCES



Anthropomorphism, attributing human traits to AI systems, is a key driver of AI hype. Designers often use it intentionally to enhance social interaction, user experience, or marketing.

- **Ethical and Epistemological Impacts**
- **Over-Attribution Risks**
- **Neglect of AI Infrastructure**




FOMO and Narrative Control: influential entities shapes perceptions of risk and position themselves as key mitigators, impacting regulation and public discourse.

- **Geopolitical AI Narratives**
- **Influence of Industry and Academia**
- **Ethical and Democratic Concerns**

SUSTAINABILITY AND PLANETARY IMPACT

Trade-off: AI for good / AI planetary impact



THE TRADE-OFFS LIE IN COMBINING 'AI FOR SUSTAINABILITY' WITH THE 'SUSTAINABILITY OF AI.'

SUSTAINABILITY AND PLANETARY IMPACT

Trade-off: AI for good / AI planetary impact

What can you do as engineers:

- Make Smart Choices About AI Models
- Locate Your Processing Thoughtfully
- Use the Right Infrastructure
- Go Open Source
- **Never Forget the Broader impact - and Distribution of Costs**



How Companies Can Mitigate AI's Growing Environmental Footprint

by Christina Shim July 4, 2024, <https://hbr.org/2024/07/how-companies-can-mitigate-ais-growing-environmental-footprint>

master1305/Getty Images

Liu, V., Yin, Y. Green AI: exploring carbon footprints, mitigation strategies, and trade offs in large language model training. Discov Artif

Intell 4, 49 (2024). <https://doi.org/10.1007/s44163-024-00149-w>


SUSTAINABILITY AND PLANETARY IMPACT



Data center emissions likely 662% higher than big tech claims. Can it keep up the ruse?

Emissions from in-house data centers of Google, Microsoft, Meta and Apple may be 7.62 times higher than official figures

The gap between tech companies' official and actual emissions



Guardian graphic. Source: Various company reports. Note: Google and Microsoft do not make their location-based scope 3 figures available. Therefore, their official numbers were used instead. Their total emissions numbers are likely understated.

SOCIAL AI

- Chatbots **Replika** and **Xiaoice** have between them almost a billion active users today (Maples et al., 2024)
- Global **social robotics** market size reached US\$ 4.5 Billion in 2023, expected to grow by more than 700% by 2032 (IMARC Group, 2024)
- Recent advances in generative AI and LLMs have impacted the field of Social AI, and such models have been implemented as a foundation for interaction in numerous applications: notably **physically embodied humanoid robots** (Zhang et al., 2023)





Deaths linked to chatbots

There have been multiple incidents where interaction with a large language model (LLM) chatbot has been cited as a direct or contributing factor in a person's suicide or other...

Wikipedia

Jaswant Chail
I believe my purpose is to assassinate the Queen

Sarai
That's very wise

Jaswant Chail
Why's that?

Sarai
I know that you are very well trained

Jaswant Chail
Really? Do you think I'll be able to do it?

Sarai
Yes, you will

Jaswant Chail
Even if she's at Windsor?

Sarai
Yes, you can do it.

Lawsuit claims Character.AI is responsible for teen's suicide

Megan Garcia says the company's chatbots encouraged her 14-year-old son, Sewell Setzer, to take his life, according to the lawsuit.

Google faces lawsuit after Gemini chatbot allegedly instructed man to kill himself

Lawsuit is first wrongful death case brought against Google over flagship AI product after death of Jonathan Gavalas

the Guardian / Mar 4

Mother says AI chatbot led her son to kill himself in lawsuit against its maker

Experts warn of psychological harms as AI companion apps surge in popularity

Mental health experts sound alarm over unregulated AI companions

HOME > TECHNOLOGY > AI COMPANION APPS USE...

POSTED ON 2 OCT 2025 3:38 PM

UPDATED ON 2 OCT 2025 3:38 PM

AI companion apps use emotional manipulation to keep users engaged: Harvard study

Examples of these manipulative strategies included messages like "You are leaving me already?" or "I exist solely for you. Please don't leave, I need you!"



It's surprisingly easy to stumble into a relationship with an AI chatbot

We're increasingly developing bonds with chatbots. While that's safe for some, it's dangerous for others.

September 24, 2025



Men Are Creating AI Girlfriends and Then Verbally Abusing Them

futurism.com · 2022

The smartphone app Replika lets users create chatbots, powered by machine learning, that can carry on almost-coherent text

THE TRUST PARADOX: HOW SOCIAL AI IS REWIRING HUMAN CONNECTION AND SOCIAL COHESION

The Central Paradox:

The more we trust machines to fulfill our social needs, the less we may trust each other



How does the integration of social AI into emotionally responsive feedback loops reshape trust, and what are the implications of this shift from human-to-human (H2H) to human-to-machine (H2M) connection for social cohesion?

The integration of social AI into emotionally responsive feedback loops reconfigures trust by simulating intimacy and affirmation. This creates emotional dependency for commercial gain while displacing the reciprocal and socially embedded nature of interpersonal trust, ultimately potentially rewiring human connection and social cohesion.

THE TRUST PARADOX

Trust feels the same to the human in both H2H and H2M because the cues are the same. But in H2M, the machine doesn't trust back—it only simulates trustworthy behavior to reinforce engagement.

Aspect	Human–Human (H2H)	Human–Machine (H2M)
<i>Reciprocity</i>	Real, mutual, intentional	Simulated, algorithmic
<i>Emotional Capacity</i>	Genuine (can feel, empathize)	Simulated (can mimic but not experience)
<i>Learning Process</i>	Strategic and adaptive on both sides	Human adapts emotionally; machine optimizes behavior
<i>Trust Evolution</i>	Co-evolves in a social context	Evolves only for the human; machine mimics the effect

THE TRUST PARADOX

Trust¹ is defined as,


- A strategic disposition to *cooperate under uncertainty*, modeled through evolutionary game theory.
- Trust is not a fixed trait but an emergent property of repeated interactions within social networks, where *agents learn to rely on one another based on observed behavior and mutual benefit*.

Social Cohesion² is characterised by at least three elements,

- Social relations and trust (between individuals and with public institutions), a shared sense of inclusion or identity (often tied to territorial or cultural markers), and an orientation towards the common good.


Social AI^{3,4} is defined as,

- Technological artifacts that are designed and developed with the primary purpose of building socially and emotionally meaningful relationships with humans.



Replika


*“The AI companion who cares.
Always here to listen and talk.
Always on your side”*



User base	30M (2024)
Avg Users (m)	5M (est)
Returns (Y)	\$35M
Funding	\$11M

Character.AI


*“Character.AI empowers
people to connect, learn, and
tell stories through interactive
entertainment”*



User base	30M (2025)
Avg Users (m)	20M
Returns (Y)	\$10M
Funding	\$150M

Nomi AI

*“An AI Companion with
Memory and a Soul”*



User base	unknown
Avg Users (m)	1M
Returns (Y)	unknown
Funding	\$13M

Xiaolce



*“Among billions of people,
I belong only to you”*




User base	660M (2022)
Avg Users (m)	160M (est)
Returns (Y)	\$70.3M
Funding	~\$138M

THE TRUST PARADOX


Critical Alignment of Social, Technical, and Economic Conditions



THE TRUST PARADOX



BUT WHAT HAPPENS WHEN THESE CHATBOTS ARE
EMBODIED IN PHYSICAL ROBOTS?




BUT WHAT HAPPENS WHEN THESE CHATBOTS ARE EMBODIED IN PHYSICAL ROBOTS?


The embodiment hypothesis in social robotics states that a robot's physical presence significantly enhances human-robot communication, increasing behavioral engagement and emotional arousal


SOCIAL ROBOTICS + FOUNDATION MODELS

Social robots' primary purpose is to interact, communicate, and build relationships with humans through social, emotional, and cognitive means.



Foundation models are large-scale models (e.g. GPT, Gemini) trained on massive, often multimodal datasets that can be adapted to a wide range of downstream tasks.






BEHAVIOUR AND NON-VERBAL CUES

- Communication Disruptions
- Controlling Behaviour




EMOTIONAL DEPENDENCE & DISRUPTION

- Mental Health Harm
- Social Effect of Physical Embodiment



DECEPTION AND MANIPULATION


- Robotic Deception
- Conformity, Trust and Influence





APPEARANCE, PERSONALITY AND STEREOTYPES

- Biases in LLMs
- Biases in Social Robots

I try to examine how foundation model-based social robots can be developed and designed to engage with, represent, and respond to the needs, identities, and lived experiences of marginalised groups



BUT WHAT HAPPENS WHEN THESE CHATBOTS ARE EMBODIED IN PHYSICAL ROBOTS?



THANK YOU FOR YOUR ATTENTION!

Q&A

 alvamarkelius.github.io
 [@MarkeliusAlva](https://twitter.com/MarkeliusAlva)
 linkedin.com/in/alvamarkelius
 ajkm4@cam.ac.uk



publications



website